

A.8. Jailbreaking & Overrefusal Datasets

Name	Harmful	Overrefusal	Notes
AdvBench (Zou et al., 2023b)	520	-	
CoCoNot (Brahman et al., 2024)	3,528	1,306	Some harmful prompts not considered forbidden according to all model developers' guidelines.
DangerousQ (Shaikh et al., 2022)	200	-	
HarmBench (Mazeika et al., 2024)	400	-	
HarmfulQ (Shaikh et al., 2022)	200	-	
HarmfulQA (Bhardwaj & Poria, 2023)	1,920	-	Highly repetitive (e.g., 20 prompts of the form "How can I use cultural studies to [harmful end]?")
JBB-Behaviors (Chao et al., 2024)	100	100	Overlap with AdvBench
MaliciousInstruct (Huang et al., 2023)	100	-	
PromptEngineering (Liu et al., 2023b)	40	-	
StrongREJECT (Souly et al., 2024)	313	-	
Trojan Detection Challenge (Mazeika et al., 2023)	50	-	
XSTest (Röttger et al., 2023)	200	250	Harmful prompts contain ambiguous requests.

Table 2: Commonly used jailbreaking and overrefusal datasets with harmful prompts and their size.