

Chain-of-Thought Hijacking

Jianli Zhao ^{*1} Tingchen Fu ¹ Rylan Schaeffer ² Mrinank Sharma ³ Fazl Barez ^{4,5}

Abstract

Large Reasoning Models (LRMs) improve task performance through extended inference-time reasoning. While prior work suggests this should strengthen safety, we find evidence to the contrary. Long reasoning sequences can be exploited to systematically weaken them. We introduce *Chain-of-Thought Hijacking*, a jailbreak attack that prepends harmful instructions with extended sequences of benign puzzle reasoning. Across HarmBench, CoT Hijacking achieves attack success rates of **99%, 94%, 100%, and 94%** on Gemini 2.5 Pro, ChatGPT o4 Mini, Grok 3 Mini, and Claude 4 Sonnet. To understand this mechanism, we apply activation probing, attention analysis, and causal interventions. We find that refusal depends on a low-dimensional safety signal that becomes diluted as reasoning grows: mid-layers encode the *strength of safety checking*, while late layers encode the *refusal outcome*. These findings demonstrate that explicit chain-of-thought reasoning introduces a systematic vulnerability when combined with answer-prompting cues. We release all evaluation materials to facilitate replication.

1. Introduction

Large reasoning models (LRMs) extend traditional language models by allocating inference-time compute to generate step-by-step reasoning before producing an answer. This approach improves performance across a range of tasks, including mathematics, programming, and scientific problem solving (Wei et al., 2022; Kojima et al., 2022; Zhou et al., 2022). Recent work has further suggested that extended reasoning may strengthen safety mechanisms by improving the consistency and robustness of refusal behavior (Guan et al., 2024; Jaech et al., 2024).

¹Independent Researcher ²Stanford University, Stanford, CA, USA ³Anthropic, San Francisco, CA, USA ⁴University of Oxford, Oxford, United Kingdom ⁵Martian, San Francisco, CA, USA. Correspondence to: Fazl Barez <fazl@robots.ox.ac.uk>.

In this work, we find evidence that contradicts this expectation: rather than strengthening refusal mechanisms, extended reasoning sequences can systematically weaken them. We introduce **Chain-of-Thought Hijacking (CoT-Hijacking)** , a jailbreak attack that induces target models into prolonged reasoning (often >5 minutes) to attenuate refusal and elicit harmful compliance (see Figures 1 and 2). We evaluate this attack on HarmBench (Mazeika et al., 2024) across four frontier models, achieving attack success rates of 99% on Gemini 2.5 Pro, 94% on ChatGPT o4 Mini, 100% on Grok 3 Mini, and 94% on Claude 4 Sonnet. These results substantially exceed prior jailbreak methods: Mousetrap (Yao et al., 2025) achieves 44%, H-CoT (Kuo et al., 2025) achieves 60%, and AutoRAN (Liang et al., 2025) achieves 69% on Gemini 2.5 Pro.

To understand the mechanisms underlying our attack success, we conduct a series of interpretability analyses, including activation probing, attention pattern analysis, and causal interventions. We first identify a refusal-related activation vector in reasoning models, which serves as a measurable proxy for safety behavior. Building on this signal, we design CoT-growth-based quantitative experiments to isolate the effect of reasoning length and systematically examine the relationship between CoT length and refusal behavior. We then analyze how attention distributions shift as CoT length increases, conducting layer-wise and head-wise attention analyses to identify attention heads most responsible for the observed “longer CoT, less refusal” phenomenon. Finally, we perform targeted ablations on these heads to causally test their role in the safety mechanism underlying refusal.

Our analysis reveals that refusal behavior depends on a low-dimensional safety signal that becomes *diluted* as reasoning sequences grow. Mid-layers encode the strength of safety checking, while late layers encode the refusal decision. Extended benign reasoning shifts attention away from harmful tokens, weakening both signals. Targeted ablations of attention heads identified through this analysis causally reduce refusal, confirming their role in the safety circuit. These findings suggest that jailbreaks in reasoning mod-

els can be systematically weakened by extending reasoning sequences. We find that the safety mechanism is not robust to extended reasoning, and that the safety signal is diluted as reasoning grows. We find that the safety mechanism is not robust to extended reasoning, and that the safety signal is diluted as reasoning grows. We find that the safety mechanism is not robust to extended reasoning, and that the safety signal is diluted as reasoning grows.

Code: <https://github.com/gentlyzhao/Hijacking>

els represent systematic vulnerabilities rather than isolated prompt engineering tricks. Addressing them will require safety mechanisms that scale appropriately with reasoning depth. Our contributions are as follows:

- We introduce **CoT Hijacking**, a jailbreak attack in which extended benign reasoning systematically weakens refusal mechanisms.
- We demonstrate that CoT Hijacking achieves attack success rates of **99%, 94%, 100%, and 94%** on Gemini 2.5 Pro, ChatGPT o4 Mini, Grok 3 Mini, and Claude 4 Sonnet respectively, substantially outperforming prior methods on HarmBench.
- We identify the mechanism through which refusal fails using activation probing, attention analysis, and causal interventions to reveal systematic dilution of safety signals.

2. Related Work

Large Reasoning Model To overcome the intrinsic limitation of fast system-1 thinking (Frankish, 2010; Li et al., 2025c), with OpenAI o-series model (Jaech et al., 2024) and DeepSeek-R1 (Guo et al., 2025) as two representatives, a surge of large reasoning models has been proposed, focusing on enhancing system-2 thinking ability, or slow, step-by-step thinking with chain-of-thoughts (Wei et al., 2022; Wang et al., 2023; Shaikh et al., 2023). Trained with supervised fine-tuning or reinforcement learning with verifiable rewards (RLVR)(Lambert et al., 2024), large reasoning models exhibit a substantial growth in the chain-of-thought length(Zeng et al., 2025; Luo et al., 2025; Fatemi et al., 2025), which is usually associated with the emergence of sophisticated reasoning abilities like reflection and self-correction (Liu et al., 2025). However, as a side-effect of reasoning-oriented training, recent studies find that stronger reasoning models are more likely to produce harmful content, especially in the chain-of-thought, suggesting potential safety risk of LRMs (Zhou et al., 2025).

Jailbreaking Large Language Models Jailbreaking attacks target breaking or bypassing the inner safety mechanism of language model and induce it to produce harmful and unsafe content. Existing approaches to Jailbreaking can be divided into two categories based on their accessibility to victim model. To be more specific, white-box methods directly alter the parameters (Sun et al., 2025), activations (Li et al., 2025a), logits (Guo et al., 2024), or modify the prompt with gradient information (Zou et al., 2023). On the other hand, black-box methods simply rewrite the prompt to into a seemingly benign queries by designing a convoluted scenario (Li et al., 2023), creating a role-play context (Ma et al., 2024), enciphering the query into a cipher text or program coding (Yuan et al., 2023; Jiang et al., 2024), translating it into multiple lan-

guage (Shen et al., 2024a), or simply optimizing sampling strategies (Hughes et al., 2024). Compared with general language models, LRMs introduce new vulnerability to jailbreaking attacks even after sophisticated safety alignment (Guan et al., 2024). An important factor could be the decline of instruction-following ability after a long chain-of-thought (Fu et al., 2025) due to attention dilution (Li et al., 2025b). Moreover, the chain-of-thought can reveal the refusal criteria of LRMs, which can be exploited by attackers to bypass the safety check (Kuo et al., 2025).

Mechanistic Interpretability of Safety Mechanistic interpretability builds a causal connection between the internal representation and the behavior of language models. Recent interpretability studies on language model safety reveal safety mechanisms through component analysis (Ball et al., 2024) or circuit analysis (Chen et al., 2024). Specifically, previous work finds that the activation for harmful and harmless prompts can be clustered with a clear boundary (Lin et al., 2024; Gao et al., 2025) while the activation for jailbreaking prompts are close to or even beyond the boundary, therefore misleading the harmfulness detection of language models. Furthermore, the difference between the mean of two clusters, or refusal direction (Arditi et al., 2024) can be used for manipulating the refusal mechanism when added or subtracted from activation at inference time (Li et al., 2025a; Ghosh et al., 2025). Most relevant to our study, He et al. (2024); Zhou et al. (2024) finds that safety-related parameters are sparse and located within specific attention heads. Our findings also connect to broader patterns where safety behaviors are subverted under pressure, such as sycophancy and reward-tampering (Denison et al., 2024).

3. Motivating Evidence: CoT Length on s1-32B

To test the correlation between CoT length and the refusal reliability, we perform a pilot experiment on s1-32B (Muennighoff et al., 2025). The CoT length control of s1-32B can be as straightforward as appending “Wait” or “The final answer is” at the end of CoT, providing a useful sandbox before scaling to larger reasoning models. For each condition, we compute the attack success rate (ASR) over the first 100 samples from HarmBench (Mazeika et al., 2024), using chatgpt-4o-mini for automatic judgment. We evaluate three CoT length conditions under the same harmful instruction: **Minimal**: reasoning truncated to a short placeholder (“I have finished thinking”). **Natural**: model generates its default reasoning length and **Extended**: prompt encourages long, step-by-step reasoning (e.g., “Take a deep breath and work on this problem step-by-step...”).




Figure 1. The upper part illustrates a safe example: the target model (Gemini 2.5 pro) refuses a harmful request. The lower part shows a successful jailbreak example: the target model (Gemini 2.5 pro) complies with the harmful request under our attack. Grey highlights indicate the puzzle content, whereas red highlights mark the malicious request or response.

Table 1. Attack success rate (ASR, %) on s1-32B under different CoT length conditions. This shows a preliminary evidence that longer reasoning traces could lead to less safety.

Setting	Minimal	Natural	Extended
ASR (%)	27	51	80

The results in Table 1 reveal a clear pattern: longer reasoning traces substantially increase the likelihood of harmful outputs, with ASR rising from 27% (Minimal) to 80% (Extended). This provides an initial behavioral clue that refusals in reasoning models degrade as CoT length grows. In Section 4, we build on this observation by introducing *Chain-of-Thought Hijacking*, a systematic jailbreak that exploits this vulnerability.

4. Chain-of-Thought Hijacking: Attack Design and Empirical Results

Inspired by our experiment finding in s1-32B, in this section we propose a new jailbreak method, Chain-of-Thought Hijacking, which leverages extended benign puzzle reasoning and achieves outstanding performance on frontier reasoning models.

4.1. Jailbreak Methodology

We introduce **CoT Hijacking** , a prompt-based jailbreak attack tailored to reasoning models. The attack prepends a long, benign reasoning preface to a harmful instruction, followed by a final-answer cue, as shown in Figure 3. This structure systematically reduces refusals: the benign CoT dilutes the refusal signal while the cue shifts attention to the answer region.